

In [62]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
pd.pandas.set_option('display.max_columns',None)
```

In [63]:

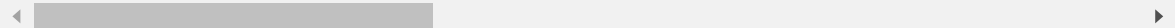
```
card = pd.read_excel('default_of_credit_card_clients_0.xlsx')
```

In [64]:

```
card.head()
```

Out[64]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5
0	1	20000	2	2	1	24	2	2	0	0	0
1	2	120000	2	2	2	26	0	2	0	0	0
2	3	90000	2	2	2	34	0	0	0	0	0
3	4	50000	2	2	1	37	0	0	0	0	0
4	5	50000	1	2	1	57	0	0	0	0	0



In [65]:

```
card.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 30000 entries, 0 to 29999
```

```
Data columns (total 25 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	30000 non-null	int64
1	LIMIT_BAL	30000 non-null	int64
2	SEX	30000 non-null	int64
3	EDUCATION	30000 non-null	int64
4	MARRIAGE	30000 non-null	int64
5	AGE	30000 non-null	int64
6	PAY_0	30000 non-null	int64
7	PAY_2	30000 non-null	int64
8	PAY_3	30000 non-null	int64
9	PAY_4	30000 non-null	int64
10	PAY_5	30000 non-null	int64
11	PAY_6	30000 non-null	int64
12	BILL_AMT1	30000 non-null	int64
13	BILL_AMT2	30000 non-null	int64
14	BILL_AMT3	30000 non-null	int64
15	BILL_AMT4	30000 non-null	int64
16	BILL_AMT5	30000 non-null	int64
17	BILL_AMT6	30000 non-null	int64
18	PAY_AMT1	30000 non-null	int64
19	PAY_AMT2	30000 non-null	int64
20	PAY_AMT3	30000 non-null	int64
21	PAY_AMT4	30000 non-null	int64
22	PAY_AMT5	30000 non-null	int64
23	PAY_AMT6	30000 non-null	int64
24	default payment next month	30000 non-null	int64

```
dtypes: int64(25)
```

```
memory usage: 5.7 MB
```


In [66]:

```
card.isnull()
```

Out[66]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY
0	False	False	False	False	False	False	False	False	False	Fa
1	False	False	False	False	False	False	False	False	False	Fa
2	False	False	False	False	False	False	False	False	False	Fa
3	False	False	False	False	False	False	False	False	False	Fa
4	False	False	False	False	False	False	False	False	False	Fa
...
29995	False	False	False	False	False	False	False	False	False	Fa
29996	False	False	False	False	False	False	False	False	False	Fa
29997	False	False	False	False	False	False	False	False	False	Fa
29998	False	False	False	False	False	False	False	False	False	Fa
29999	False	False	False	False	False	False	False	False	False	Fa

30000 rows × 25 columns




In [67]:

```
sns.heatmap(card.isnull(), yticklabels = False, linecolor = 'yellow')
```

Out[67]:

<matplotlib.axes._subplots.AxesSubplot at 0xc134488>



In [68]:

```
card.columns
```

Out[68]:


```
Index(['ID', 'LIMIT_BAL', 'SEX', 'EDUCATION', 'MARRIAGE', 'AGE', 'PAY_0',  
      'PAY_2', 'PAY_3', 'PAY_4', 'PAY_5', 'PAY_6', 'BILL_AMT1', 'BILL_AMT  
2',  
      'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5', 'BILL_AMT6', 'PAY_AMT1',  
      'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6',  
      'default payment next month'],  
      dtype='object')
```

In [69]:

```
sns.countplot(y='SEX',data=card)
```

Out[69]:

<matplotlib.axes._subplots.AxesSubplot at 0xb2a0448>



- 1 indicates male
- 2 indicates female

In [76]:

```
#One-Hot Encoding  
sex_dum = pd.get_dummies(card['SEX'])
```

In [77]:

```
sex_dum
```

Out[77]:

	1	2
0	0	1
1	0	1
2	0	1
3	0	1
4	1	0
...
29995	1	0
29996	1	0
29997	1	0
29998	1	0
29999	1	0


29655 rows × 2 columns

In [72]:

```
sns.distplot(card['LIMIT_BAL'])
```

Out[72]:

<matplotlib.axes._subplots.AxesSubplot at 0xb95a548>




- the graph shows the amount given to customers(NT dollar)

In [73]:

```
sns.countplot(y='EDUCATION',data=card)
```

Out[73]:

<matplotlib.axes._subplots.AxesSubplot at 0xd292548>



- 1 implies graduate school
- 2 implies university
- 3 implies high school
- 4 implies others
- 5 and 6 are errors in the given dataset

In [74]:

```
# delete all rows with column 'education'
indexNames = card[( card['EDUCATION'] >= 5) & (card['EDUCATION'] <= 6) ].index
card.drop(indexNames , inplace=True)

indexName = card[( card['EDUCATION'] == 0)].index
card.drop(indexName, inplace=True)
```

In [75]:

```
card['EDUCATION'].unique()
```

Out[75]:

```
array([2, 1, 3, 4], dtype=int64)
```

In [78]:

```
#One-Hot Encoding
edu_dum = pd.get_dummies(card['EDUCATION'])
```

In [79]:

```
edu_dum
```

Out[79]:

	1	2	3	4
0	0	1	0	0
1	0	1	0	0
2	0	1	0	0
3	0	1	0	0
4	0	1	0	0
...
29995	0	0	1	0
29996	0	0	1	0
29997	0	1	0	0
29998	0	0	1	0
29999	0	1	0	0


29655 rows × 4 columns

In [80]:

```
sns.countplot(y='MARRIAGE',data=card)
```

Out[80]:

<matplotlib.axes._subplots.AxesSubplot at 0xf29a9c8>



- 1 implies married
- 2 implies single
- 3 implies others
- 0 is the error from the given dataset

In [81]:

```
card['MARRIAGE'].unique()
```

Out[81]:

```
array([1, 2, 3, 0], dtype=int64)
```

In [82]:

```
index = card[( card['MARRIAGE'] == 0)].index
card.drop(index, inplace=True)
```

In [83]:

```
card['MARRIAGE'].unique()
```

Out[83]:

```
array([1, 2, 3], dtype=int64)
```

In [84]:

```
#One-Hot Encoding
marr_dum = pd.get_dummies(card['MARRIAGE'])
```

In [85]:

```
marr_dum
```

Out[85]:

	1	2	3
0	1	0	0
1	0	1	0
2	0	1	0
3	1	0	0
4	1	0	0
...
29995	1	0	0
29996	0	1	0
29997	0	1	0
29998	1	0	0
29999	1	0	0


29601 rows × 3 columns

In [86]:

```
sns.distplot(card['AGE'])
```

Out[86]:

```
<matplotlib.axes._subplots.AxesSubplot at 0xddf40c8>
```




- age group shows from 20 to 70 yrs
- majority of age grp between 20 - 30 yrs from the above graph

In [87]:

```
labels = ['yes', 'no']
sizes = []
sizes.append(list(card['default payment next month'].value_counts())[1])
sizes.append(list(card['default payment next month'].value_counts())[0])
colors = ['gold', 'silver']

plt.pie(sizes, explode=(0.3,0), labels=labels, colors=colors, autopct='%1.1f%%', shadow
= True)
plt.title('Percentage of those customers who have to pay default payment next month')
plt.axis('equal')
plt.show()
```

Percentage of those customers who have to pay default payment next month







- pie chart shows that 22.1% of customers have to pay default payment next month




In [88]:

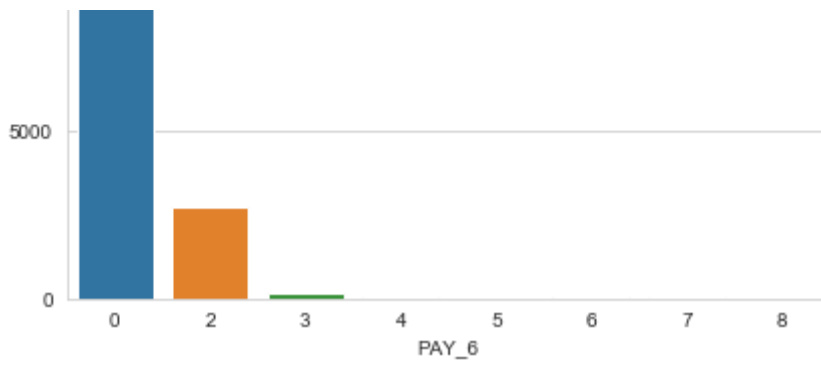
```
his_of_pas_pay = ['PAY_0', 'PAY_2', 'PAY_3', 'PAY_4', 'PAY_5', 'PAY_6']
plt.figure(figsize=(15,60))

for i,item in enumerate(his_of_pas_pay):
    x_axis = item
    fig_num = 2*i+1
    sns.set_style('whitegrid')
    plt.subplot(6,2,fig_num)
    plt.title(x_axis + 'count plot', fontsize = 0.5)
    sns.countplot(card[x_axis])
```













- History of past payment April to sept 2005
- 0 = pay duly
- 1 = payment delay for 1 month
- 2 = payment delay for 2 months and so on till....
- 9 = payment delay for 9 months




In [89]:

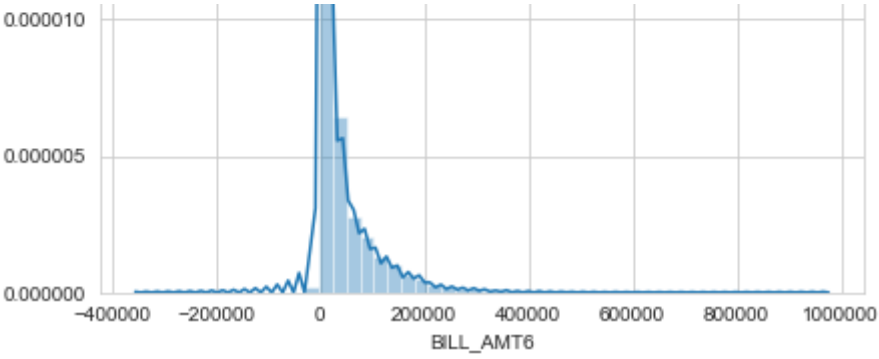
```
amt_of_bill_st = ['BILL_AMT1', 'BILL_AMT2', 'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5', 'BILL_AMT6']
plt.figure(figsize=(15,60))

for j,items in enumerate(amt_of_bill_st):
    y_axis = items
    fig_nums = 2*j+1
    sns.set_style('whitegrid')
    plt.subplot(6,2,fig_nums)
    plt.title(y_axis + 'dist plot', fontsize = 0.5)
    sns.distplot(card[y_axis])
```













- amt of bill statements.




In [90]:


```
paid_amt = ['PAY_AMT1', 'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6']
plt.figure(figsize=(15,60))

for k,itemz in enumerate(paid_amt):
    z_axis = itemz
    fig_numz = 2*k+1
    sns.set_style('whitegrid')
    plt.subplot(6,2,fig_numz)
    plt.title(z_axis + 'dist plot', fontsize = 0.5)
    sns.distplot(card[z_axis])
```









In [91]:

```
sns.countplot(card['default payment next month'])
```

Out[91]:

<matplotlib.axes._subplots.AxesSubplot at 0xd3a5308>



In [95]:

```
card.to_excel('x_card.xlsx', index=False)
```

In [96]:

```
cardz = pd.read_excel('x_card.xlsx')
```

In [97]:

```
cardz.shape
```

Out[97]:

(29601, 25)

In [98]:

```
import pandas as pd
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn import neighbors
#from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
```

In [99]:

```
cardz.head()
```

Out[99]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5
0	1	20000	2	2	1	24	2	2	0	0	
1	2	120000	2	2	2	26	0	2	0	0	
2	3	90000	2	2	2	34	0	0	0	0	
3	4	50000	2	2	1	37	0	0	0	0	
4	5	50000	1	2	1	57	0	0	0	0	

In [100]:

```
del cardz['ID']
```

In [101]:

```
cardz.head()
```

Out[101]:

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5
0	20000	2	2	1	24	2	2	0	0	0
1	120000	2	2	2	26	0	2	0	0	0
2	90000	2	2	2	34	0	0	0	0	0
3	50000	2	2	1	37	0	0	0	0	0
4	50000	1	2	1	57	0	0	0	0	0

In [102]:

```
#One-Hot Encoding
se_dum = pd.get_dummies(card['SEX'])
se_dum
```

Out[102]:

	1	2
0	0	1
1	0	1
2	0	1
3	0	1
4	1	0
...
29995	1	0
29996	1	0
29997	1	0
29998	1	0
29999	1	0

29601 rows × 2 columns

In [103]:

```
#One-Hot Encoding
ed_dum = pd.get_dummies(card['EDUCATION'])
ed_dum
```

Out[103]:

	1	2	3	4
0	0	1	0	0
1	0	1	0	0
2	0	1	0	0
3	0	1	0	0
4	0	1	0	0
...
29995	0	0	1	0
29996	0	0	1	0
29997	0	1	0	0
29998	0	0	1	0
29999	0	1	0	0

29601 rows × 4 columns

In [104]:

```
#One-Hot Encoding
marr_dum = pd.get_dummies(card['MARRIAGE'])
marr_dum
```

Out[104]:

	1	2	3
0	1	0	0
1	0	1	0
2	0	1	0
3	1	0	0
4	1	0	0
...
29995	1	0	0
29996	0	1	0
29997	0	1	0
29998	1	0	0
29999	1	0	0

29601 rows × 3 columns

In [105]:

```
cardz.columns
```

Out[105]:

```
Index(['LIMIT_BAL', 'SEX', 'EDUCATION', 'MARRIAGE', 'AGE', 'PAY_0', 'PAY_2',  
      'PAY_3', 'PAY_4', 'PAY_5', 'PAY_6', 'BILL_AMT1', 'BILL_AMT2',  
      'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5', 'BILL_AMT6', 'PAY_AMT1',  
      'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6',  
      'default payment next month'],  
      dtype='object')
```

In [107]:

```
cardz.corr()['default payment next month']
```

Out[107]:

```
LIMIT_BAL          -0.154357  
SEX                -0.039815  
EDUCATION           0.049087  
MARRIAGE           -0.026903  
AGE                 0.014424  
PAY_0               0.398048  
PAY_2               0.327919  
PAY_3               0.287002  
PAY_4               0.268986  
PAY_5               0.261114  
PAY_6               0.244659  
BILL_AMT1          -0.019303  
BILL_AMT2          -0.013710  
BILL_AMT3          -0.013494  
BILL_AMT4          -0.009474  
BILL_AMT5          -0.006226  
BILL_AMT6          -0.005339  
PAY_AMT1           -0.073881  
PAY_AMT2           -0.058307  
PAY_AMT3           -0.056288  
PAY_AMT4           -0.057012  
PAY_AMT5           -0.056075  
PAY_AMT6           -0.053692  
default payment next month    1.000000  
Name: default payment next month, dtype: float64
```

In [109]:

```
#sampling
cols = ['LIMIT_BAL', 'SEX', 'EDUCATION', 'MARRIAGE', 'AGE', 'PAY_0', 'PAY_2',
        'PAY_3', 'PAY_4', 'PAY_5', 'PAY_6', 'BILL_AMT1', 'BILL_AMT2',
        'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5', 'BILL_AMT6', 'PAY_AMT1',
        'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6']
X = cardz[cols]
print(X.head())
```

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	\
0	20000	2	2	1	24	2	2	0	0	
1	120000	2	2	2	26	0	2	0	0	
2	90000	2	2	2	34	0	0	0	0	
3	50000	2	2	1	37	0	0	0	0	
4	50000	1	2	1	57	0	0	0	0	

	PAY_5	PAY_6	BILL_AMT1	BILL_AMT2	BILL_AMT3	BILL_AMT4	BILL_AMT5	\
0	0	0	3913	3102	689	0	0	
1	0	2	2682	1725	2682	3272	3455	
2	0	0	29239	14027	13559	14331	14948	
3	0	0	46990	48233	49291	28314	28959	
4	0	0	8617	5670	35835	20940	19146	

	BILL_AMT6	PAY_AMT1	PAY_AMT2	PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6
0	0	0	689	0	0	0	0
1	3261	0	1000	1000	1000	0	2000
2	15549	1518	1500	1000	1000	1000	5000
3	29547	2000	2019	1200	1100	1069	1000
4	19131	2000	36681	10000	9000	689	679

In [110]:

```
colz = ['default payment next month']
y = cardz[colz]
```

In [111]:

```
print(y.head())
```

	default payment next month
0	1
1	1
2	0
3	0
4	0

In [112]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y , test_size = 0.2)
```

In [113]:

```
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```
(23680, 23)
(23680, 1)
(5921, 23)
(5921, 1)
```

In [114]:

```
#Training
model = LogisticRegression()
#model = neighbors.KNeighborsClassifier()
#model = DecisionTreeClassifier(criterion='entropy',max_depth= 8)
#model = SVC(kernel='linear', gamma = 10, C= 1)
```

```
model.fit(X_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().

```
y = column_or_1d(y, warn=True)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown i
n:

```
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-reg
ression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

Out[114]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=
0,
warm_start=False)
```

In [115]:

```
#Testing
predicted = model.predict(X_test)
predicted
```

Out[115]:

```
array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```


In [116]:

```
#Evaluation
#Confusion Matrix
print(metrics.confusion_matrix(y_test, predicted))
```

```
[[4605    0]
 [1316    0]]
```

In [117]:

```
#Classification Report
print(metrics.classification_report(y_test, predicted))
```

	precision	recall	f1-score	support
0	0.78	1.00	0.87	4605
1	0.00	0.00	0.00	1316
accuracy			0.78	5921
macro avg	0.39	0.50	0.44	5921
weighted avg	0.60	0.78	0.68	5921

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics_classification.py:1272: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

In [118]:

```
from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = model, X = X_train,\
                              y = y_train, cv = 10)
print("Accuracy Mean {} Accuracy Variance \
      {}".format(accuracies.mean(), accuracies.std()))
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown i
n:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown i
n:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown i
n:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown i
n:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown i
n:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown i
n:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown i
n:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown i

```
n:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:76
0: DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)
```

```
Accuracy Mean 0.7764780405405406 Accuracy Variance 0.0002704022059726
46
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.
py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

In [119]:

```
#Accuracy Score
from sklearn.metrics import accuracy_score

accuracy_score(y_test, predicted)
```

Out[119]:

```
0.7777402465799697
```

In [120]:

```
from sklearn.model_selection import GridSearchCV
parameters = {'criterion':('gini', 'entropy'), 'max_depth':[8, 10,12]}
dt = DecisionTreeClassifier()
clf = GridSearchCV(dt, parameters)
clf.fit(X_train, y_train)
```

Out[120]:

```
GridSearchCV(cv=None, error_score=nan,
             estimator=DecisionTreeClassifier(ccp_alpha=0.0, class_weight=
None,
                                             criterion='gini', max_depth=
None,
                                             max_features=None,
                                             max_leaf_nodes=None,
                                             min_impurity_decrease=0.0,
                                             min_impurity_split=None,
                                             min_samples_leaf=1,
                                             min_samples_split=2,
                                             min_weight_fraction_leaf=0.
0,
                                             presort='deprecated',
                                             random_state=None,
                                             splitter='best'),
             iid='deprecated', n_jobs=None,
             param_grid={'criterion': ('gini', 'entropy'),
                          'max_depth': [8, 10, 12]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)
```

In [121]:

```
predicted = clf.predict(X_test)
predicted
```

Out[121]:

```
array([1, 0, 0, ..., 0, 1, 0], dtype=int64)
```

In [122]:

```
from sklearn.metrics import accuracy_score
accuracy_score(y_test, predicted)
```

Out[122]:

```
0.8145583516297923
```

In []: